

Jakub Háva
jakub@h2o.ai

Sparkling Water 2.0

H2O & Booking.com, Amsterdam,
April 6, 2017

Spark⁺ H₂O

SPARKLING
WATER

Who am I

- Finishing high-performance cluster monitoring tool for JVM based languages (JNI, JVMTI, instrumentation)
- Finishing Master's at Charles Uni in Prague
- Software engineer at H2O.ai - Core Sparkling Water
- Climbing & Tea lover, (doesn't mean I don't like beer!)

Distributed Sparkling Team

- Michal - Mt. View, CA
- Kuba - Prague, CZ
- Mateusz - Tokyo, JP
- Vlad - Mt. View, CA

H₂O+Spark =
Sparkling
Water

Sparkling Water

- Transparent integration of H2O with Spark ecosystem - MLlib and H2O side-by-side
- Transparent use of H2O data structures and algorithms with Spark API
- Platform for building Smarter Applications
- Excels in existing Spark workflows requiring advanced Machine Learning algorithms

Functionality missing in H2O can be replaced by Spark and vice versa

Benefits



- Additional algorithms
 - NLP
- Powerful data munging
- ML Pipelines
- Advanced algorithms
 - speed v. accuracy
 - advanced parameters
- Fully distributed and parallelised
- Graphical environment
- R/Python interface

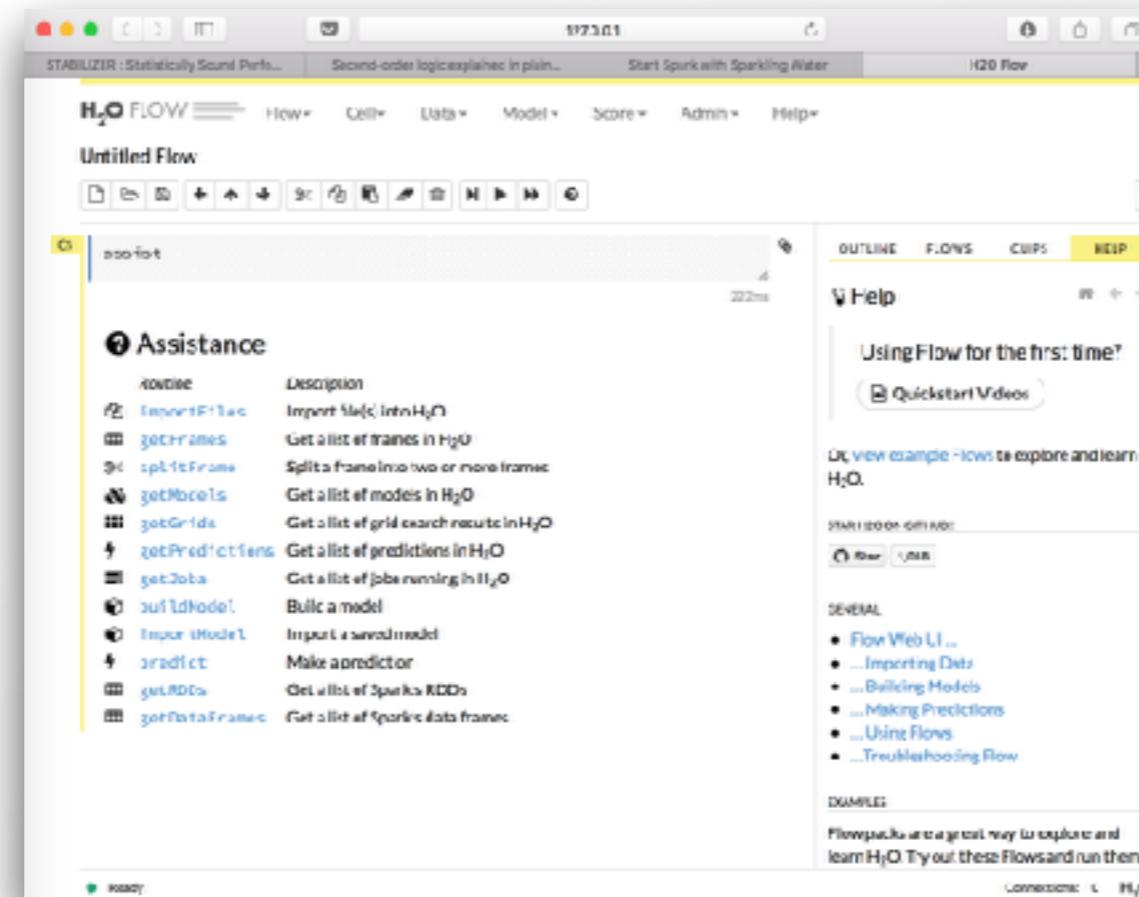
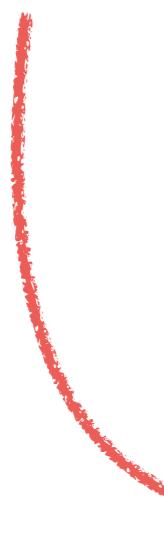
How to use Sparkling Water?

Start spark with Sparkling Water

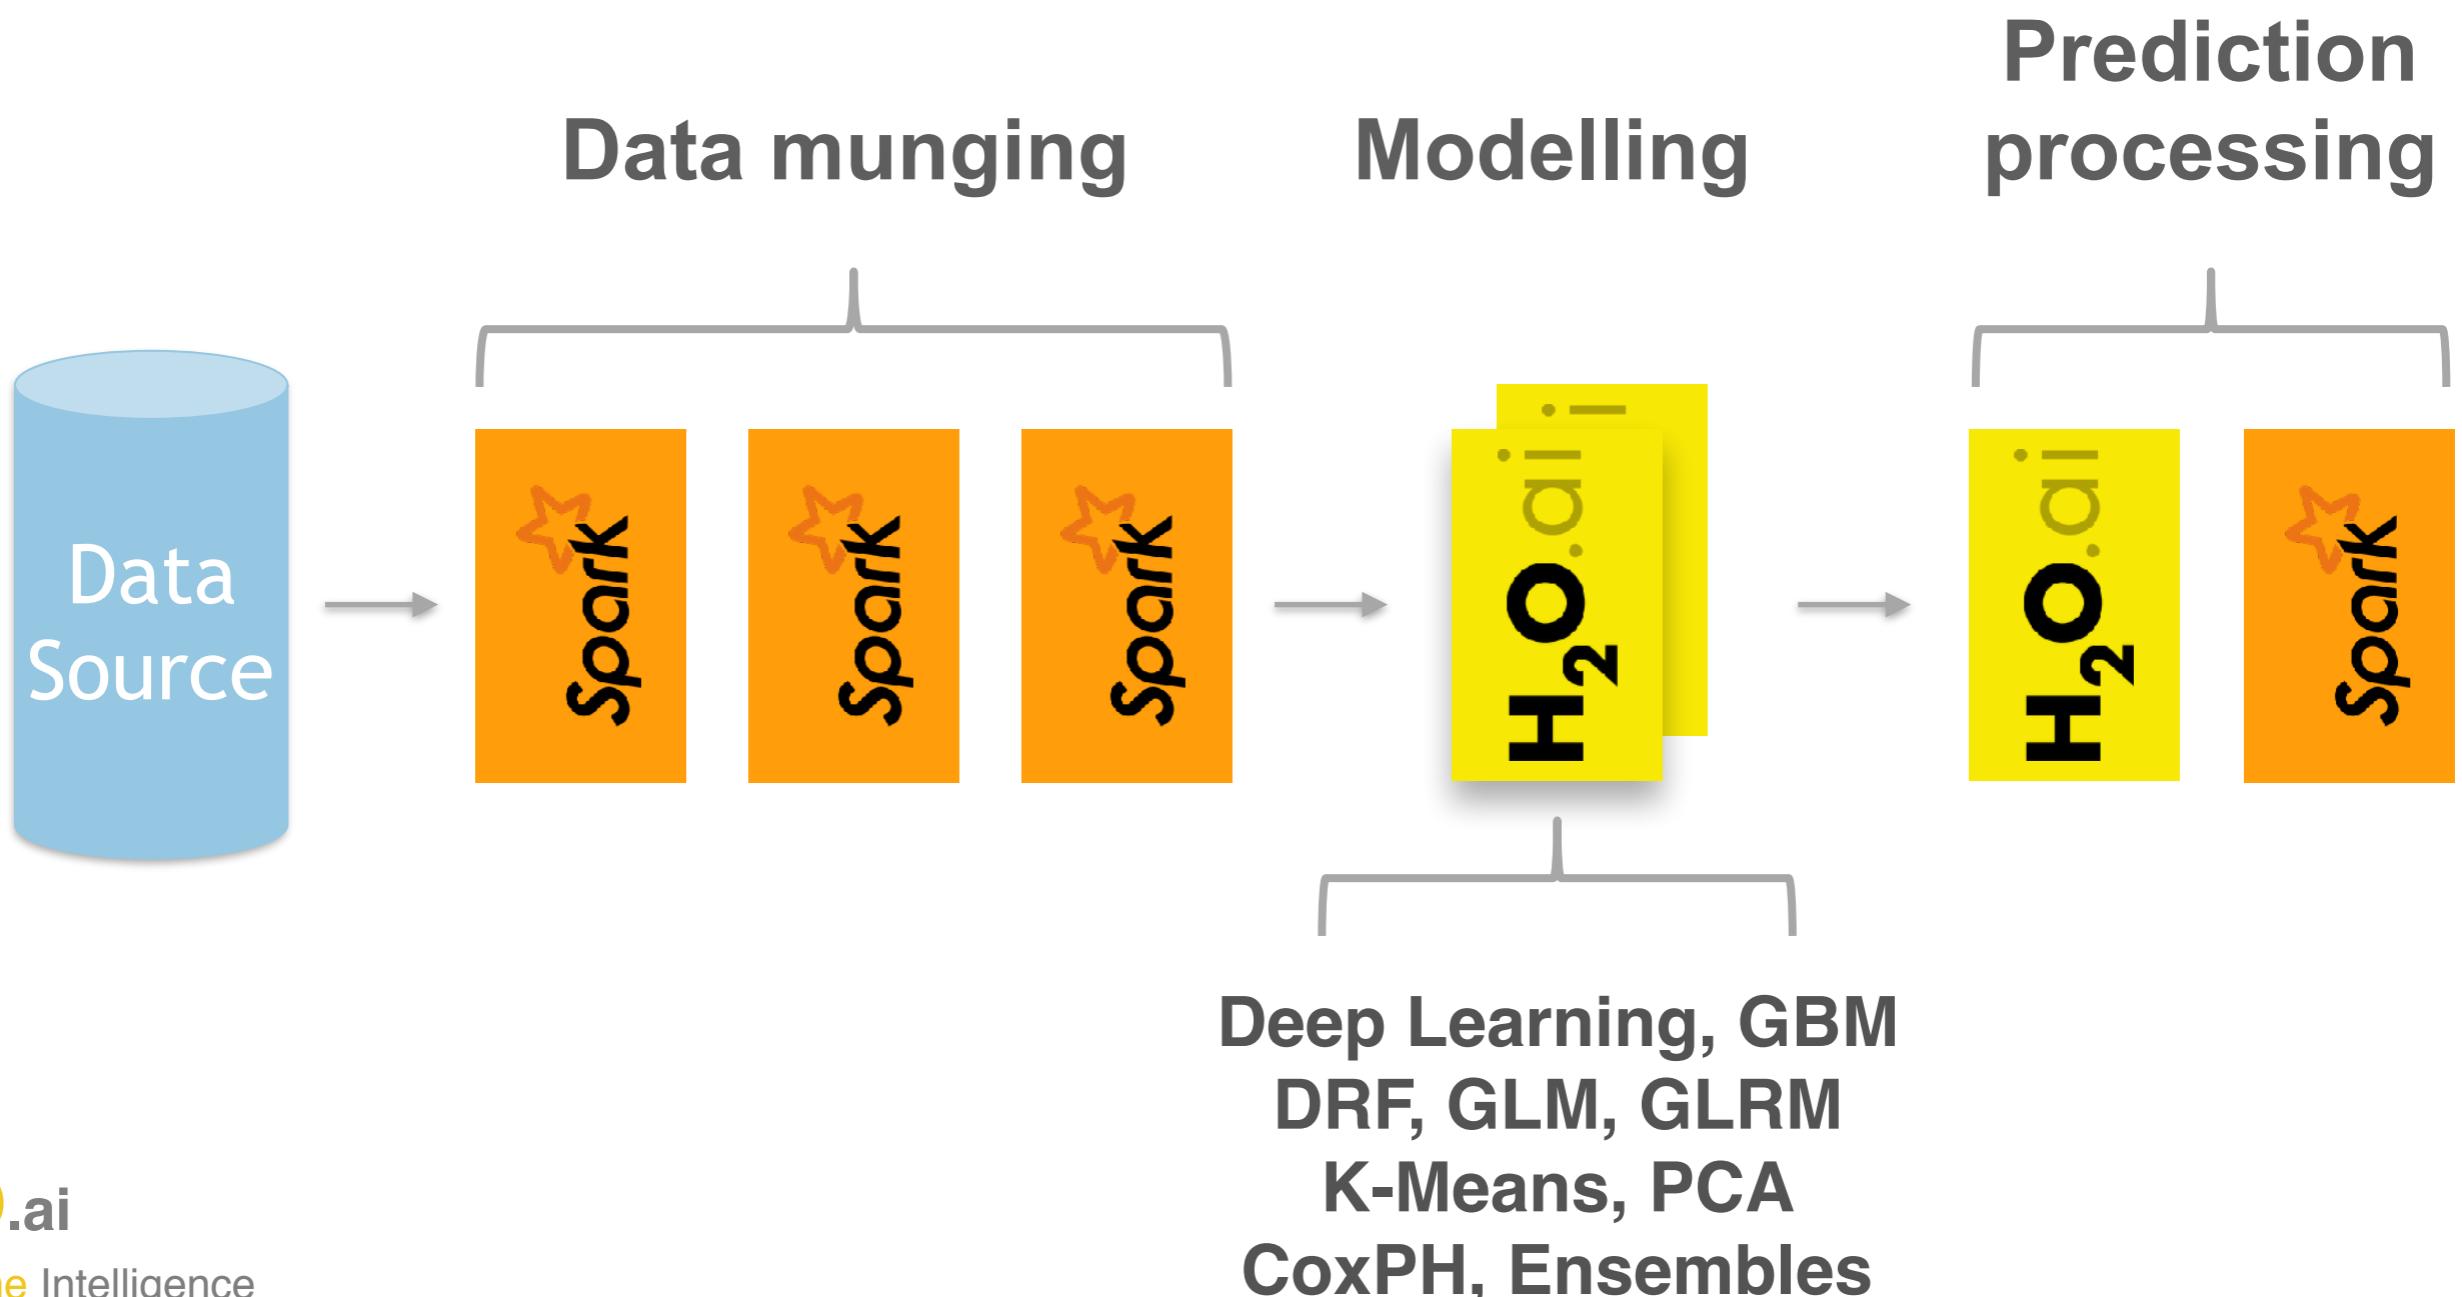
start.sh

```
1 $SPARK_HOME/bin/spark-submit \
2   --class water.SparklingWaterDriver \
3   --packages ai.h2o:sparkling-water-examples_2.10:1.6.3 \
4   --executor-memory=6g \
5   --driver-class-path scalastyle.jar /dev/null
```

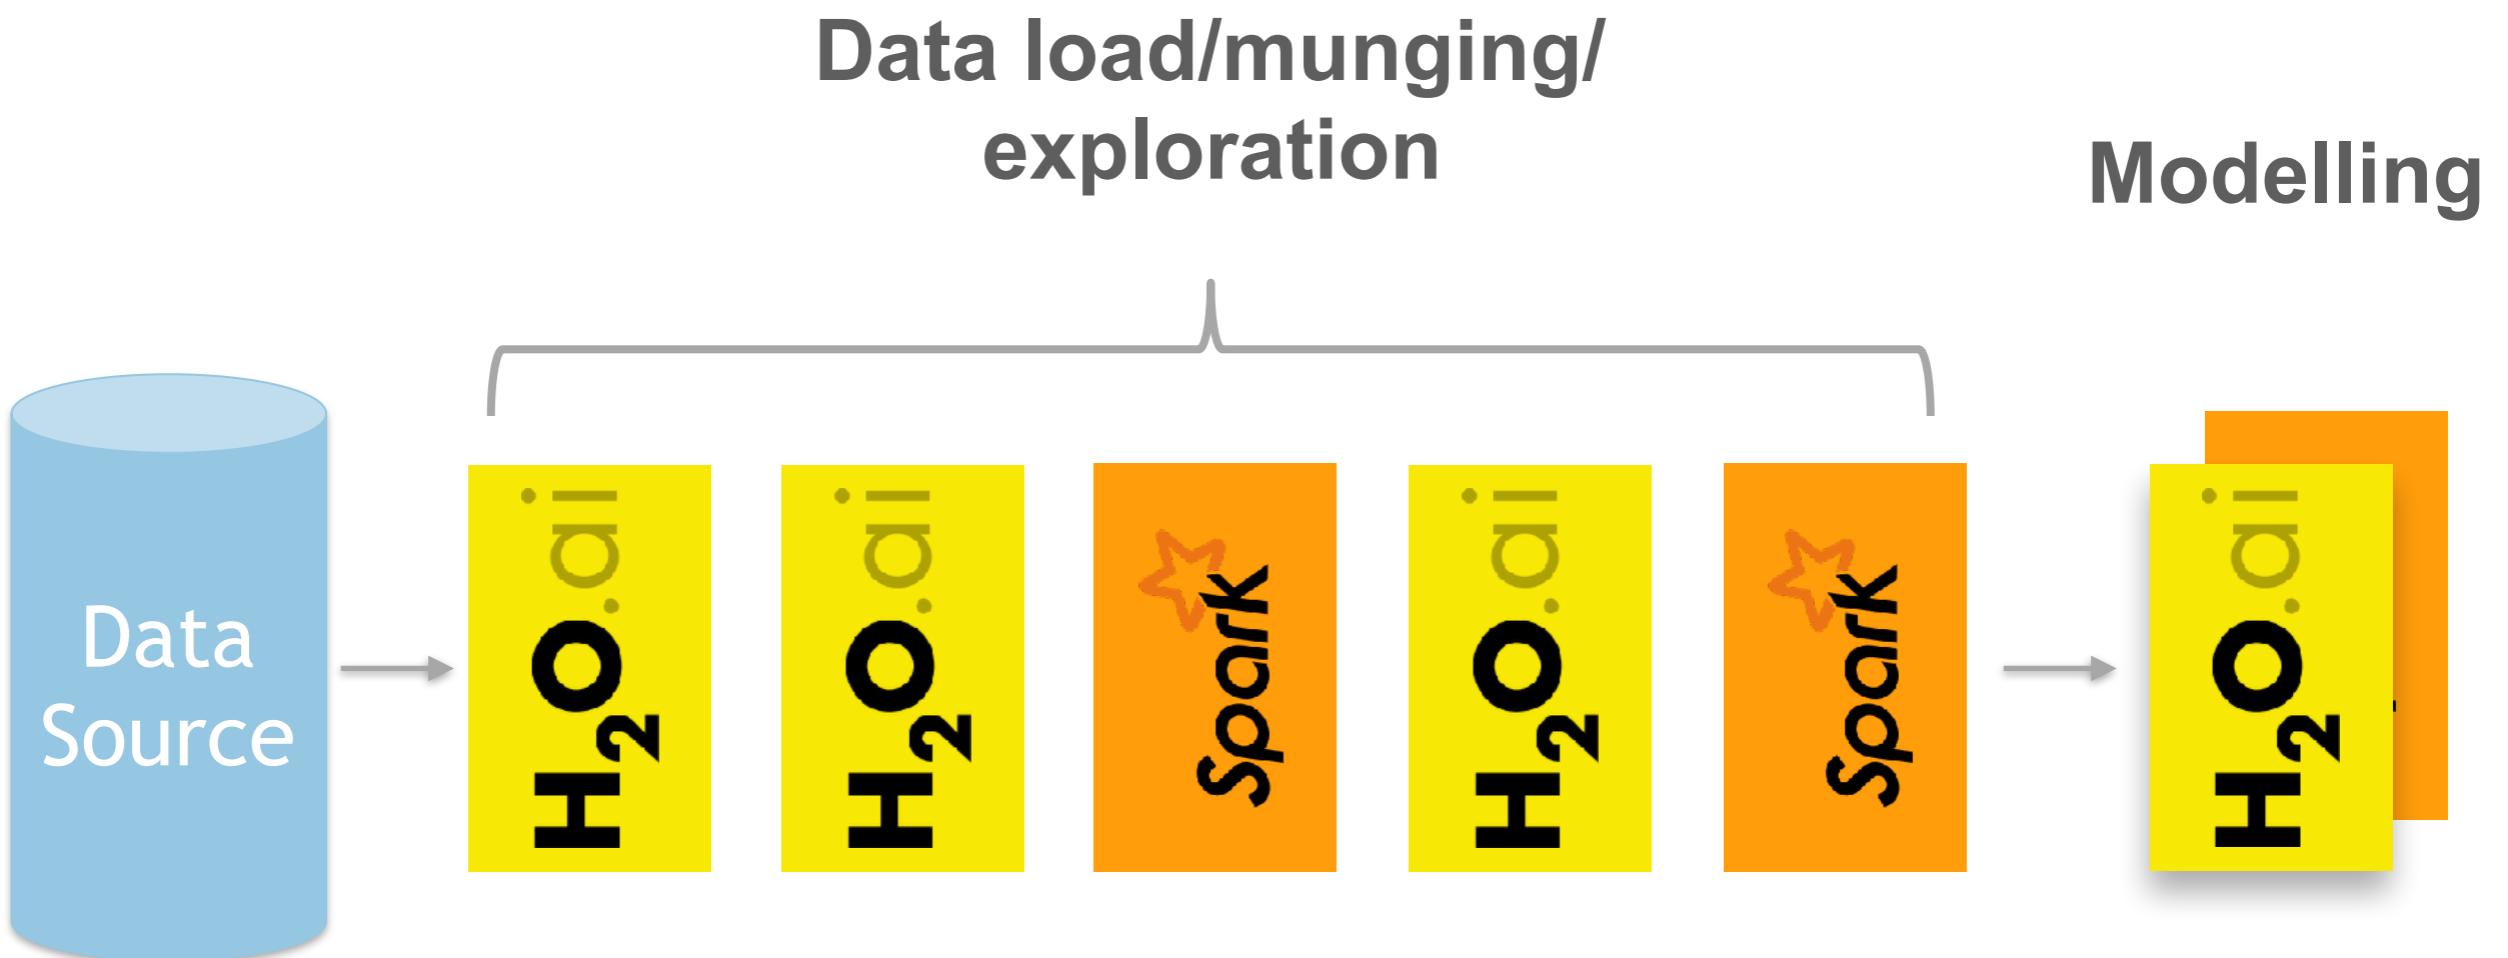
Raw



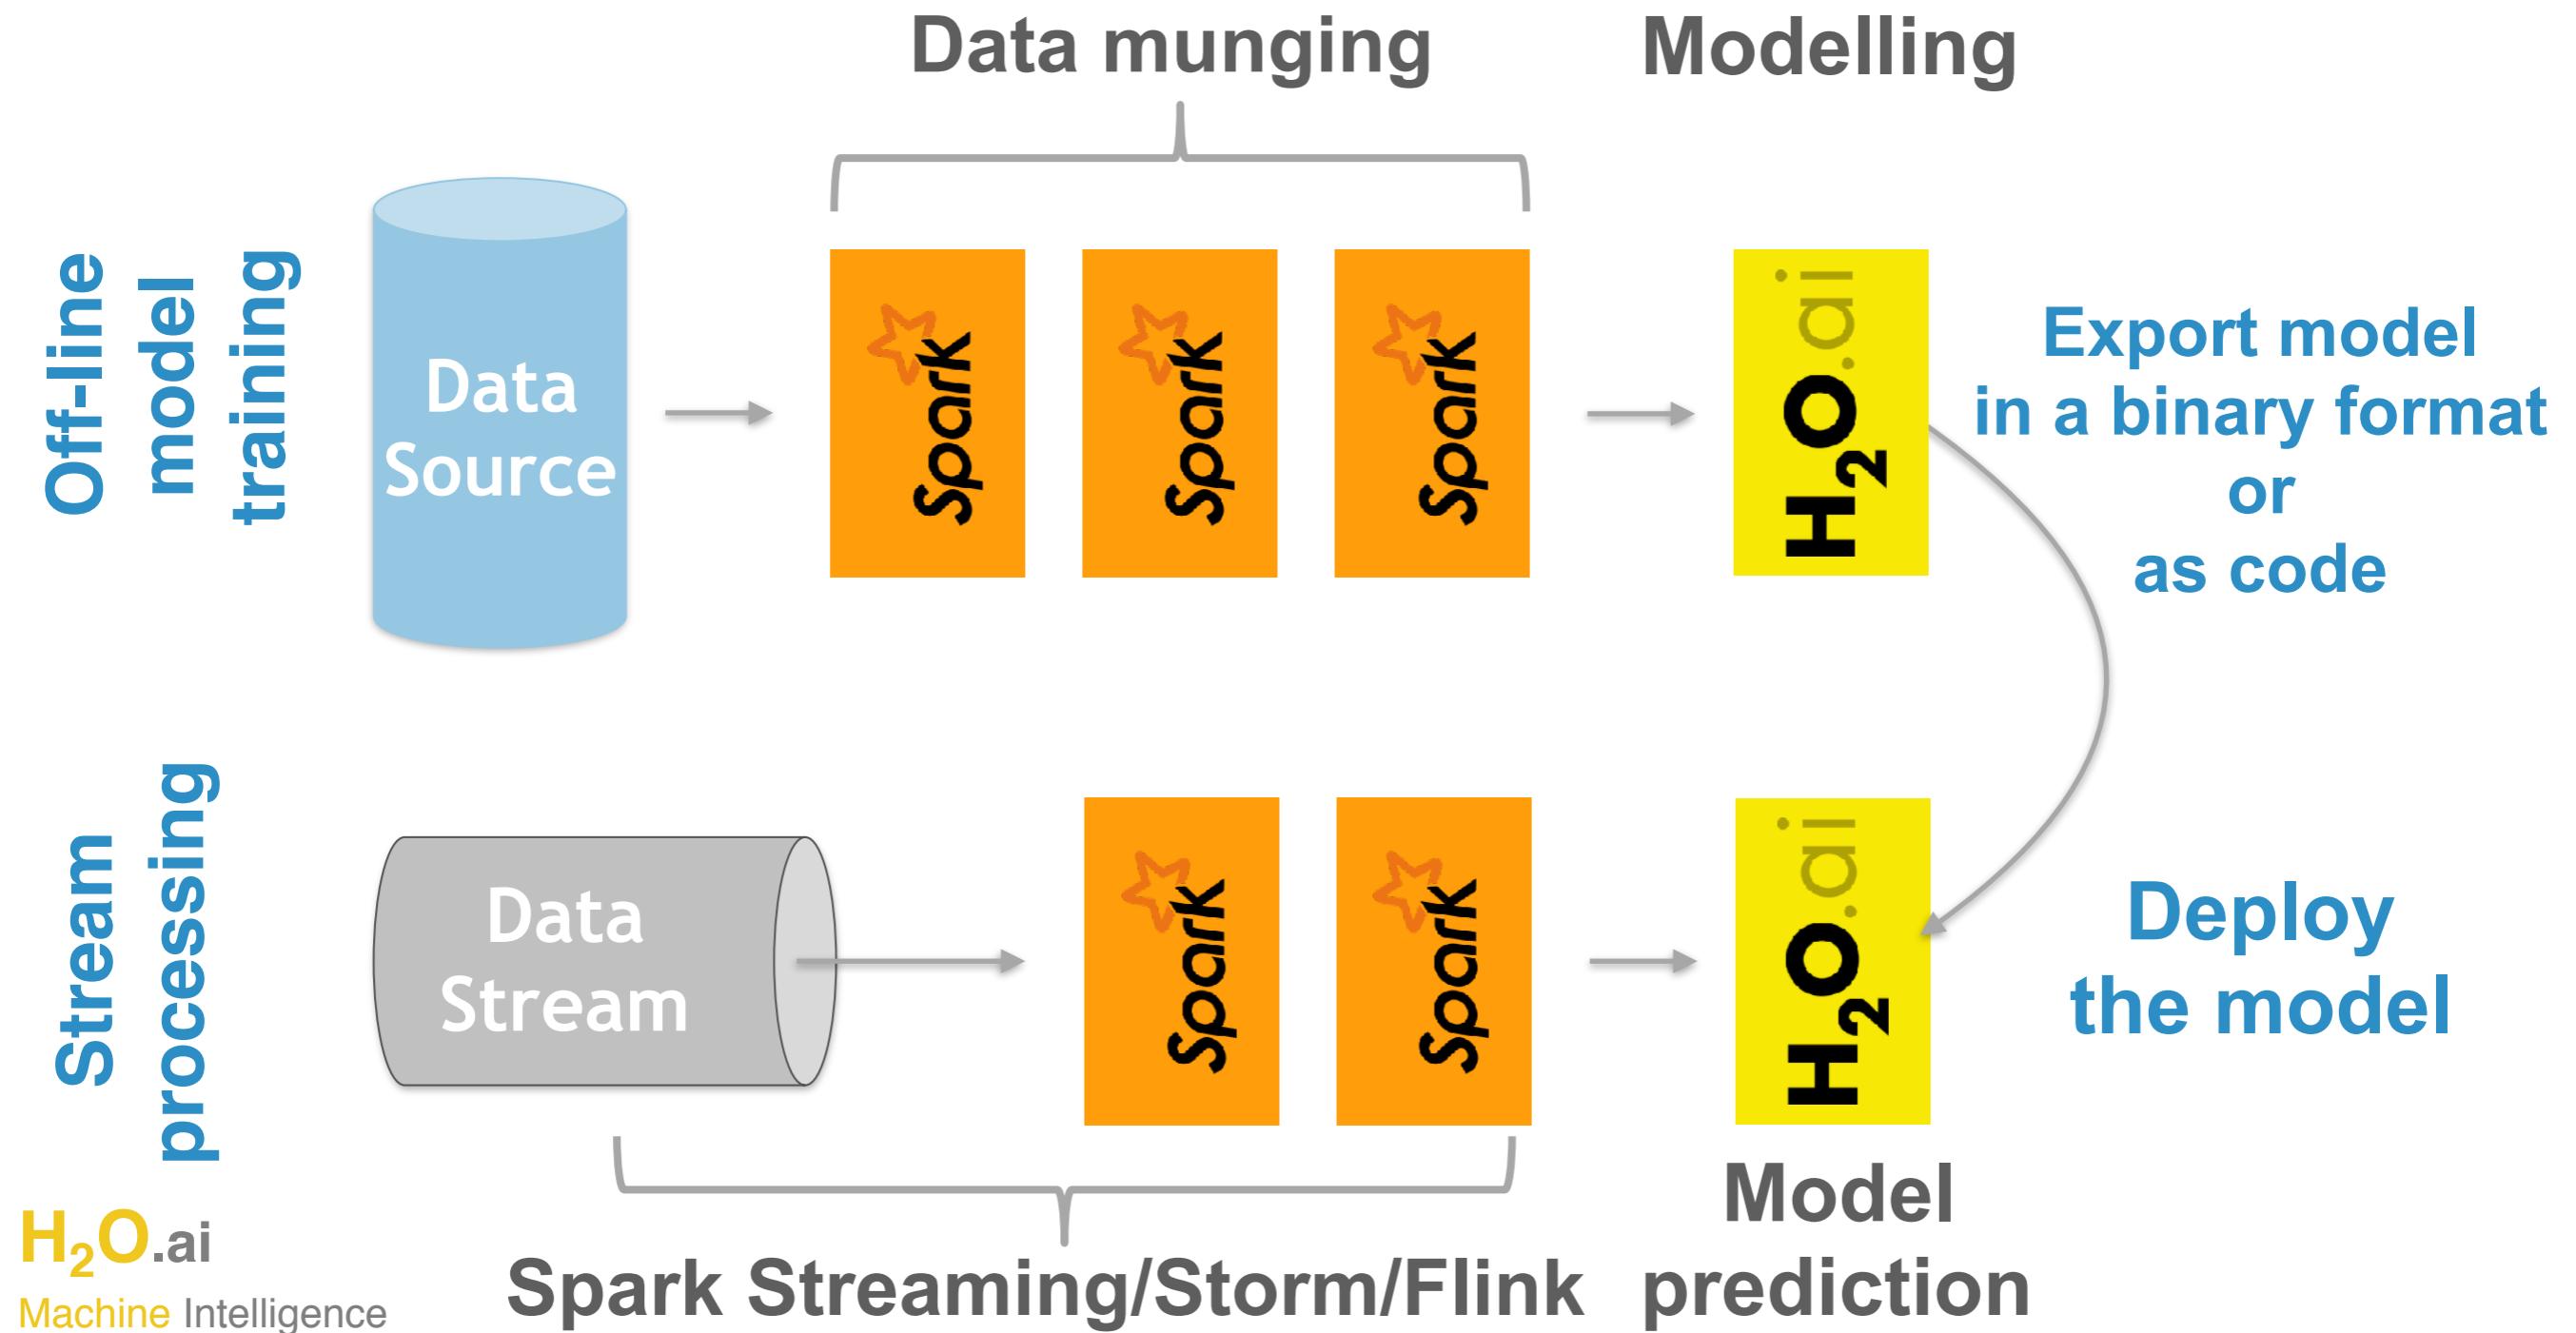
Model Building



Data Munging



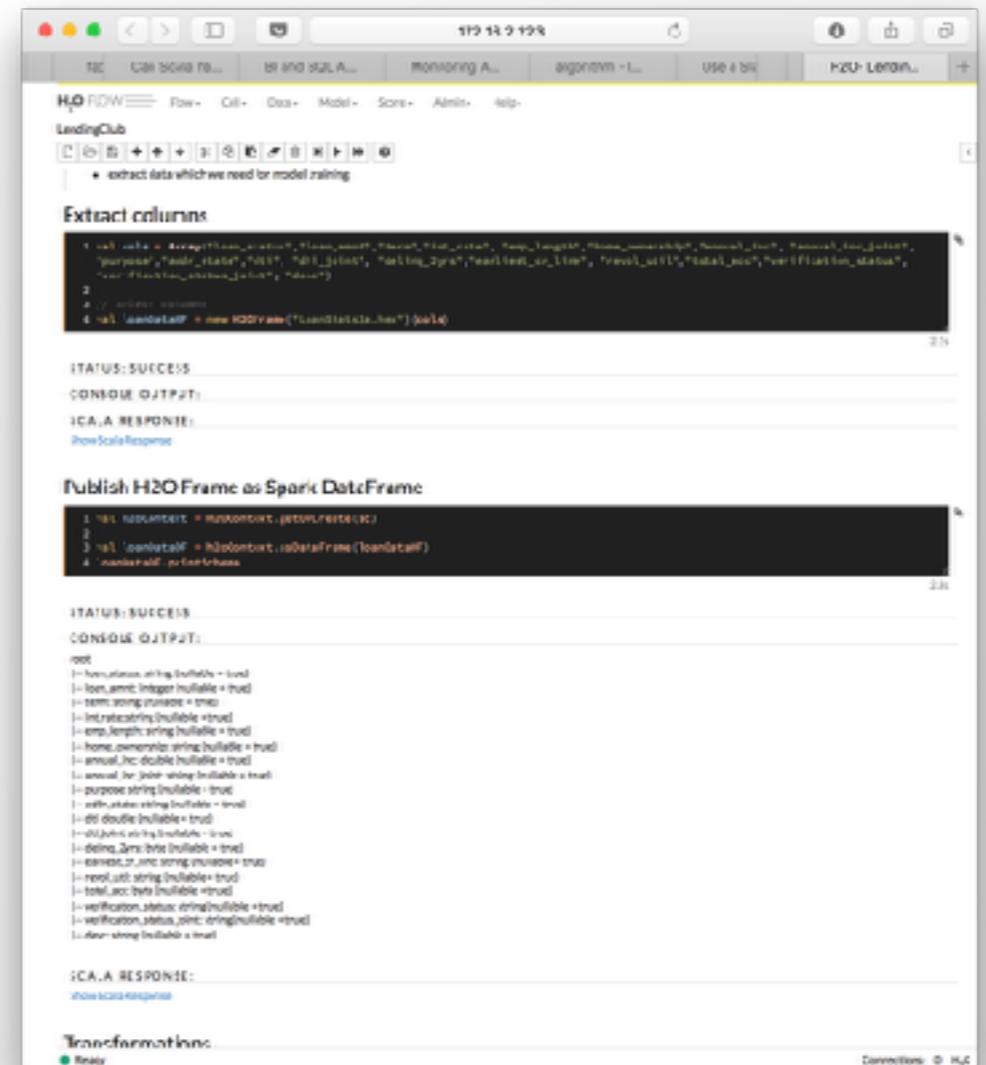
Stream processing



Features Overview

Scala code in H2O Flow

- New type of cell
 - Access Spark from Flow UI
 - Experimenting made easy



H2O Frame as Spark's Datasource

- Use native Spark API to load and save data
- Spark can optimise the queries when loading data from H2O Frame
- Use of Spark query optimiser

Machine learning pipelines

- Wrap our algorithms as Transformers and Estimators
- Support for embedding them into Spark ML Pipelines
- Can serialise fitted/unfitted pipelines
- Unified API => Arguments are set in the same way for Spark and H2O Models

MLlib Algorithms in Flow UI

- Can examine them in H2O Flow
- Can generate POJO out of them

PySparkling made easy

- PySparkling now in PyPi
- Contains all H2O and Sparkling Water dependencies, no need to worry about them
- Just add in on your Python path and that's it

And others!

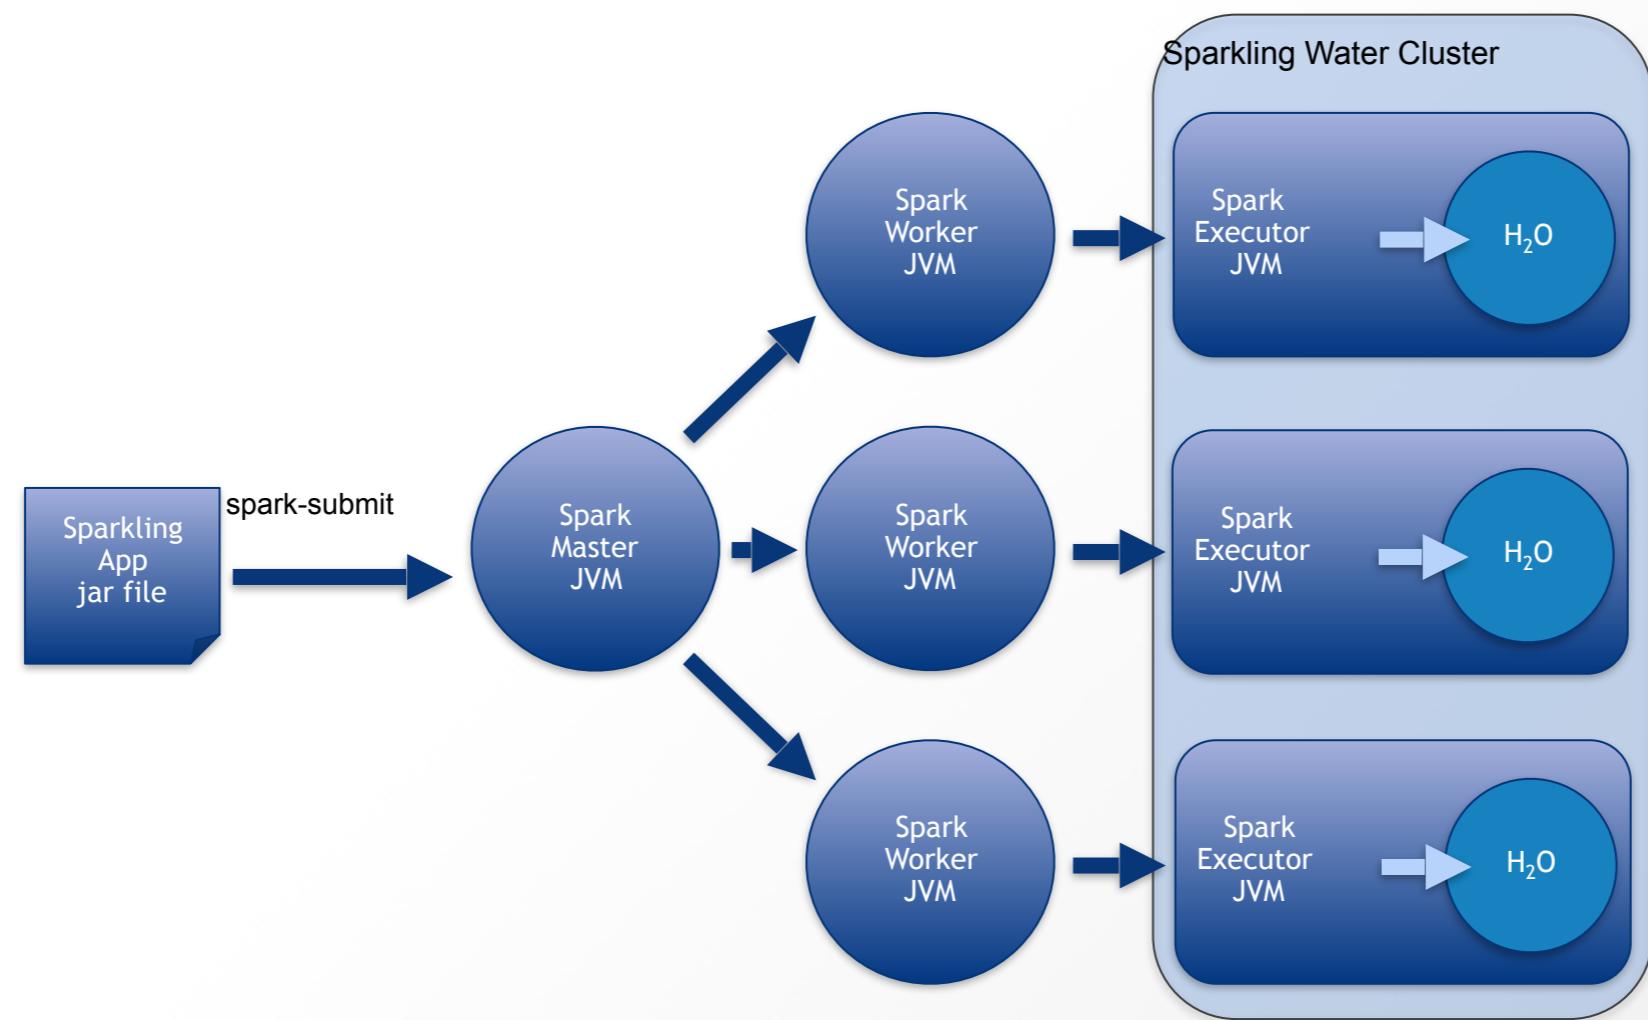
- Support for Datasets
- RSparkling
- Zeppelin notebook support
- Integration with TensorFlow (via DeepWater)
- Support for high-cardinality fast joins
- Secure Communication - SSL
- Bug fixes..

Coming features

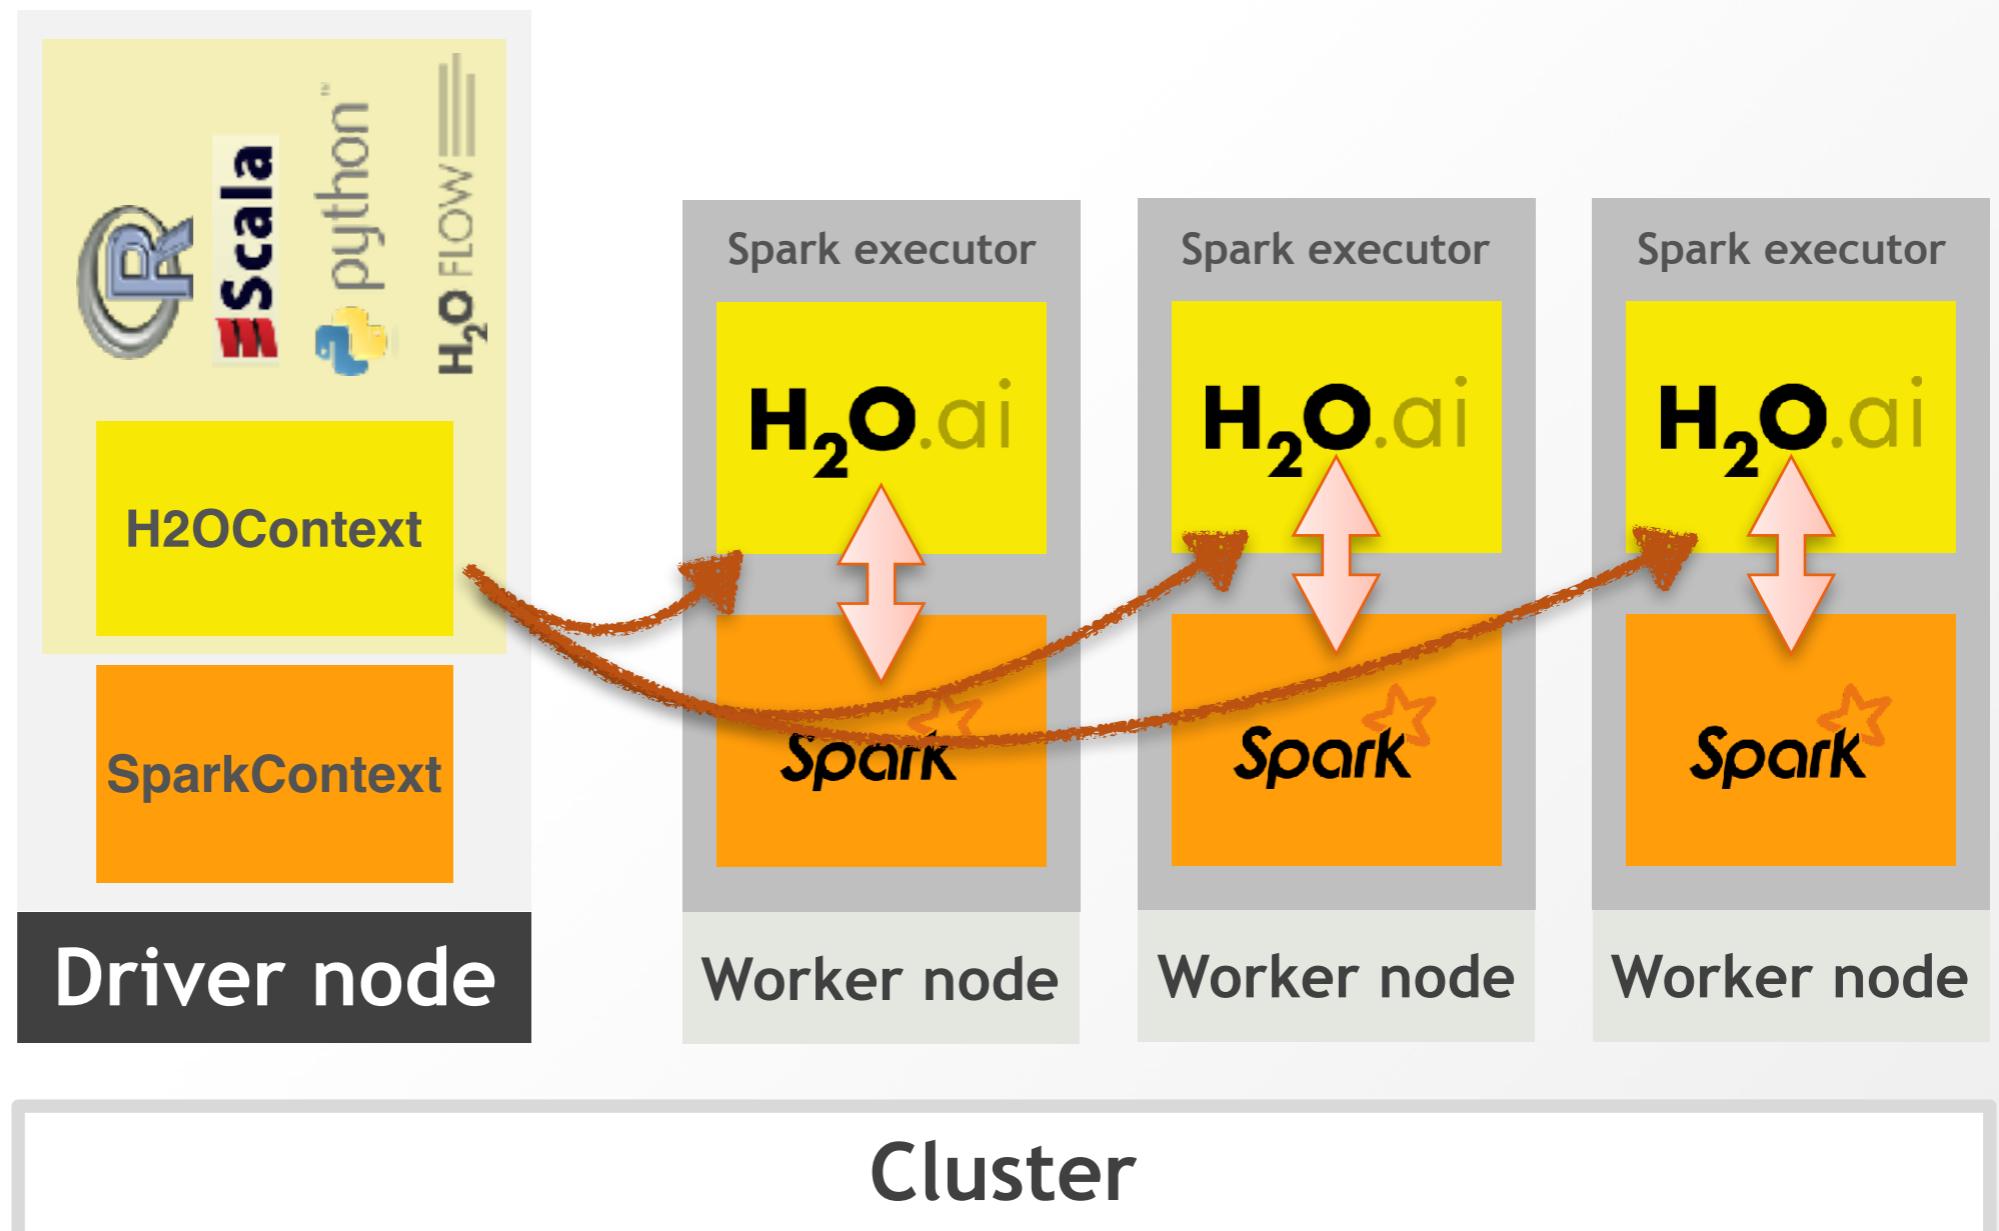
- Support for more MLlib algorithms in Flow
- Python cell in the H2O Flow
- Integration with Steam
- ...

Internal Backend

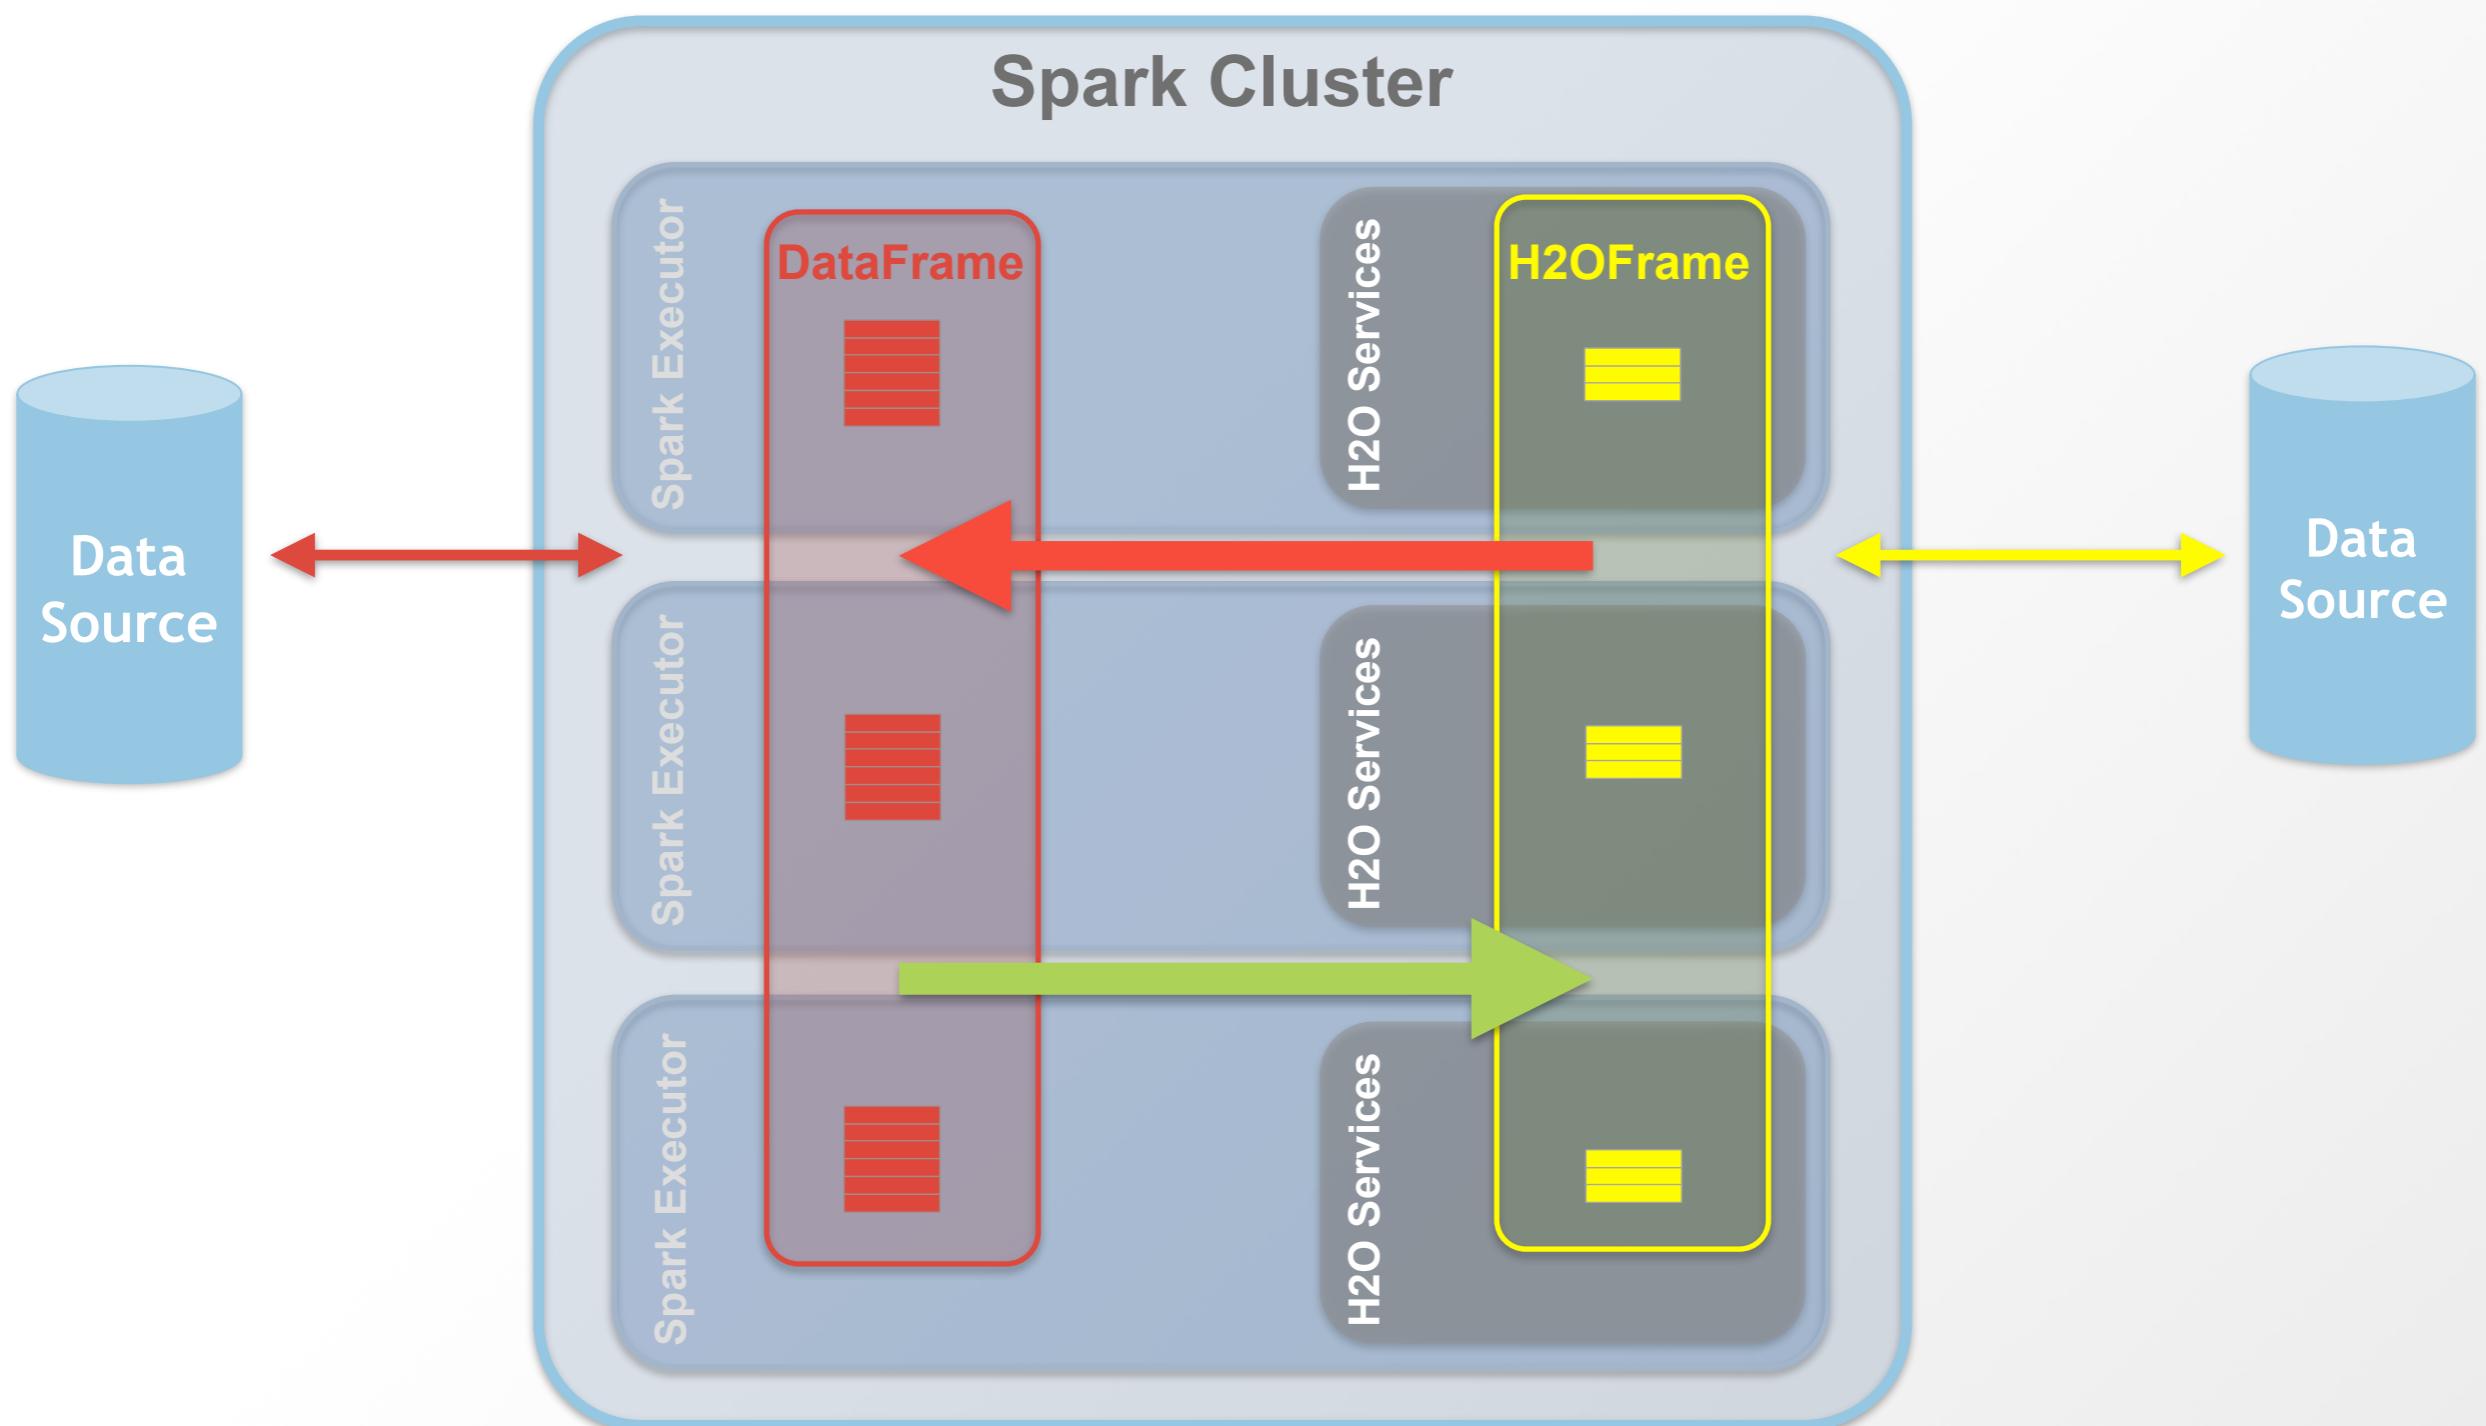
Sparkling Water Internal Backend



Internal Backend



Internal Backend



Pros & Cons

- Advantages
 - Easy to configure
 - Faster (no need to send data to different cluster)
- Disadvantages
 - Spark kills or joins new executor => H2O goes down
 - No way how to discover all executors

When to use

- Have a few Spark Executors
- Stable environment - not many Spark jobs need to be repeated
- Quick jobs

When not To Use

- Unstable environment
- Big number of executors
- Big data, long running tasks which often fails in the middle

How to use

```
from pysparkling import *

conf = H2OConf(sc)
    .set_internal_cluster_mode()
hc = H2OContext.getOrCreate(sc, conf)
```

External Backend

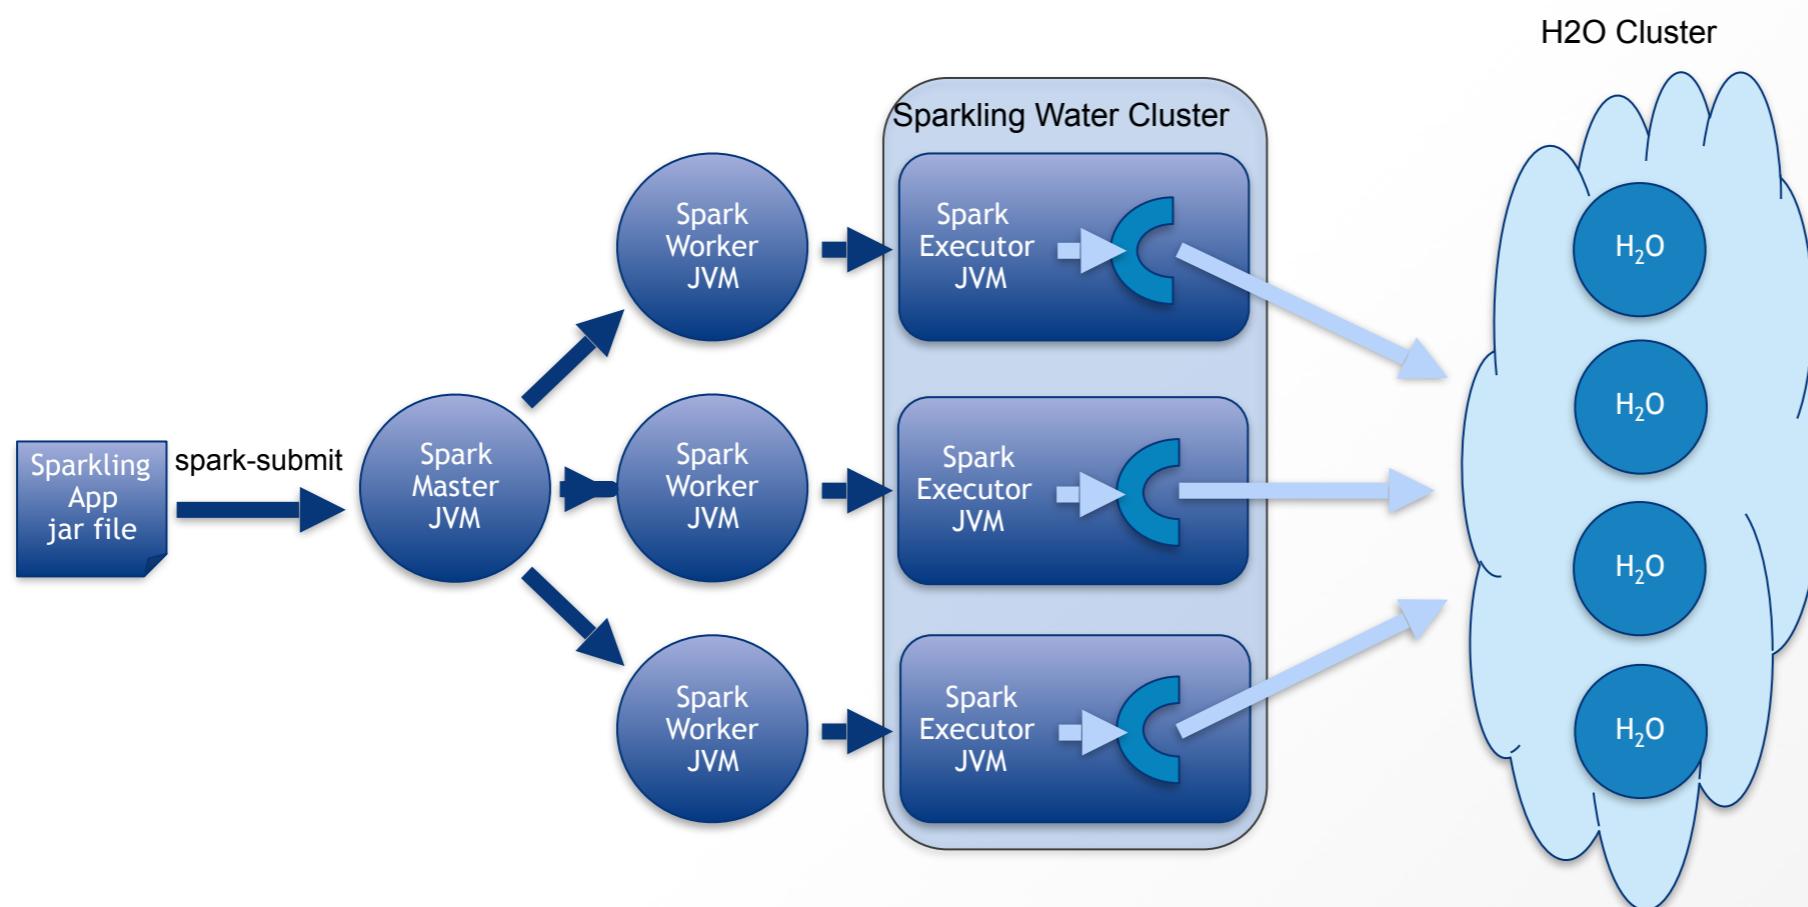
Overview

- Sparkling Water is using external H2O cluster instead of starting H2O in each executor
- Spark executors can come and go and H2O won't be affected
- Start H2O cluster on YARN automatically

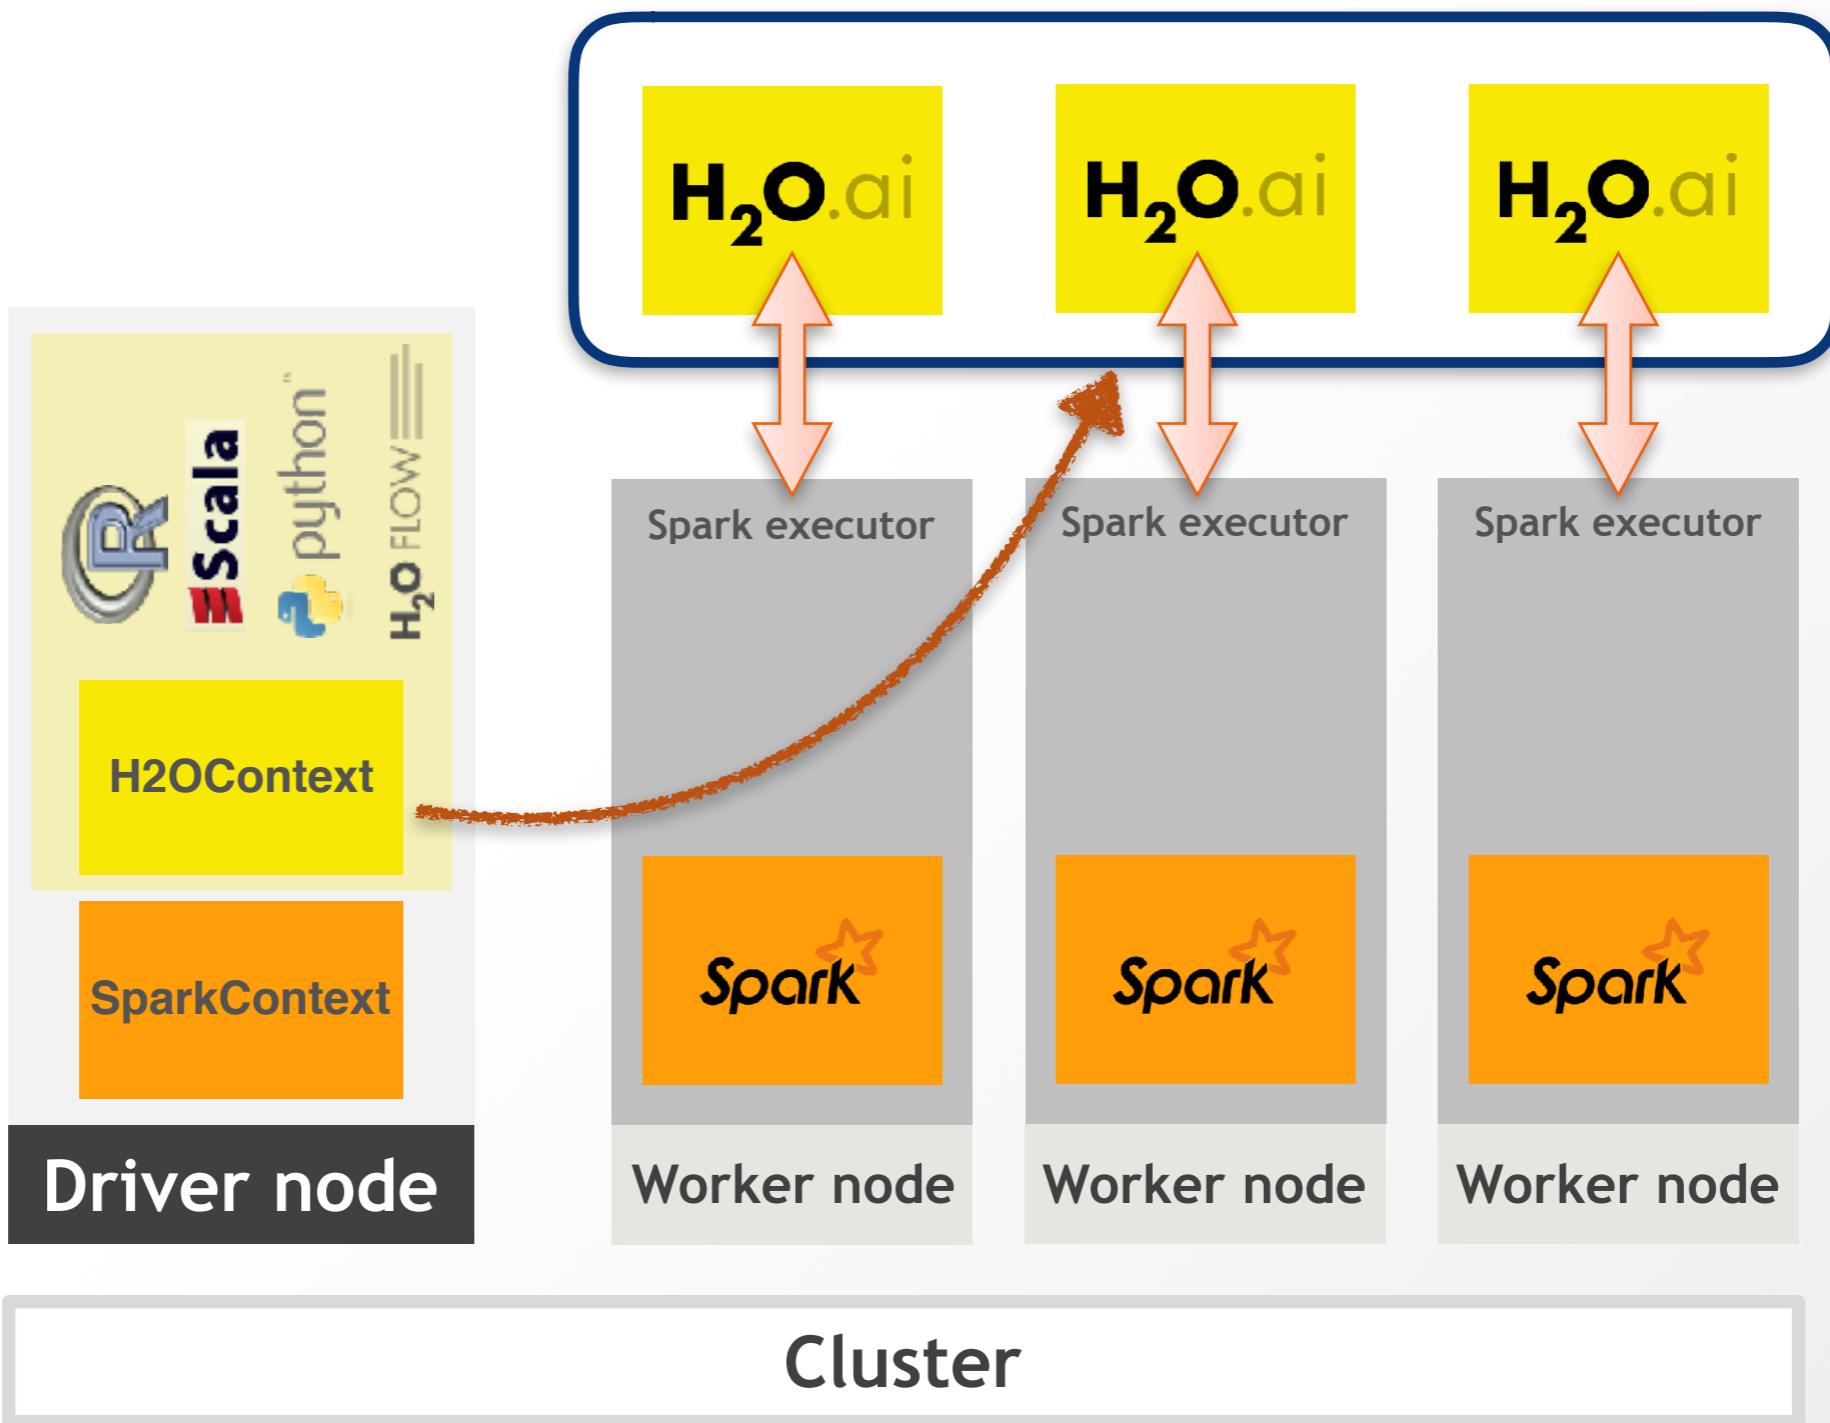
Separation Approach

- Separating Spark and H2O
 - But preserving same API:
`val h2oContext = H2OContext.getOrCreate("http://h2o:54321/")`
- Spark and H2O can be submitted as Yarn job and controlled in separation
 - But H2O still needs non-elastic environment (H2O itself does not implement HA yet)

Sparkling Water External Backend



External Backend



Pros & Cons

- **Advantages**
 - H2O does not crash when Spark executor goes down
 - Better resource management since resources can be planned per tool
- **Disadvantages**
 - Transfer overhead between Spark and H2O processes
 - under measurement with cooperation of a customer

Extended H2O Jar

- External cluster requires H2O jar file to be extended by sparkling water classes
- `./gradlew extendJar -PdownloadH2O=cdh5.8`
- We are working on making this even easier so the extended Jars will be available online

Modes

- **Auto Start mode**
 - Start h2o cluster automatically on YARN
- **Manual Start Mode**
 - User is responsible for starting the cluster manually

Manual Cluster Start Mode

```
from pysparkling import *

conf = H2OConf(sc)
    .set_external_cluster_mode()
    .use_manual_cluster_start()
    .set_cloud_name("test")
    .set_h2o_driver_path("path to h2o driver")

hc = H2OContext.getOrCreate(sc, conf)
```

Auto Cluster Start Mode

```
from pysparkling import *

conf = H2OConf(sc)
    .set_external_cluster_mode()
    .use_auto_cluster_start()
    .set_h2o_driver_path("path_to_extended_driver")
    .set_num_of_external_h2o_nodes(4)
    .set_mapper_xmx("2G")
    .set_yarn_queue("abc")

hc = H2OContext.getOrCreate(sc, conf)
```

Multi-network Environment

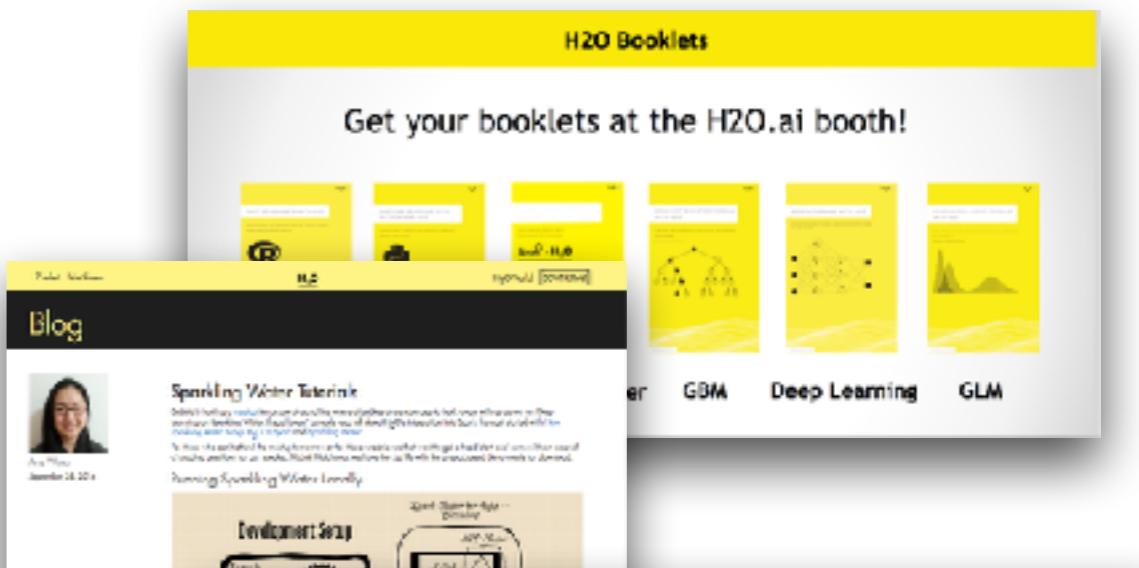
- Node with Spark driver & H2O client is connected to more networks
- Spark driver choose IP in different network then the external H2O cloud
- We need to manually configure H2O client IP to be on the same network as rest of the H2O cloud

Demo Time

More info

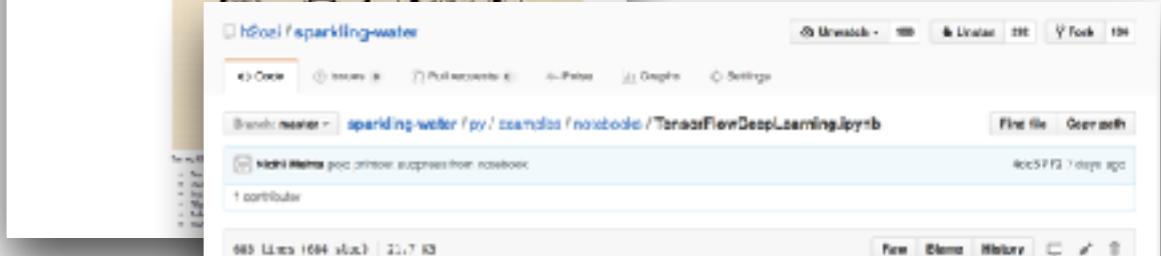
Checkout **H2O.ai** Training Books

<http://h2o.ai/resources>



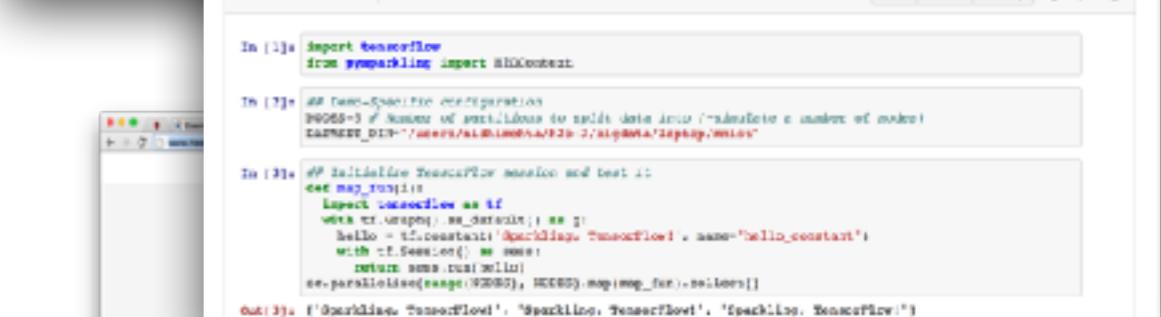
Checkout **H2O.ai** Blog

<http://h2o.ai/blog/>



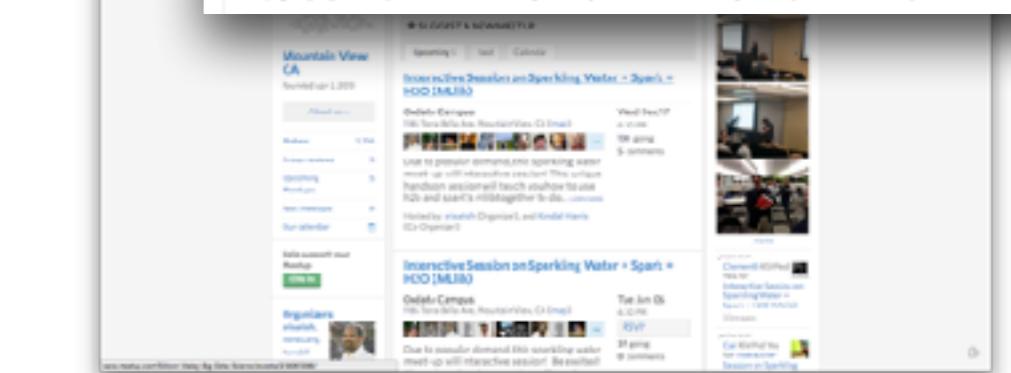
Checkout **H2O.ai** Youtube Channel

<https://www.youtube.com/user/0xdata>



Checkout GitHub

<https://github.com/h2oai/sparkling-water>



Thank you!

Sparkling Water is
open-source
ML application platform
combining
power of Spark and H2O

Learn more at h2o.ai
Follow us at [@h2oai](https://twitter.com/h2oai)

