

Overview of Machine Learning and H2O.ai

Patrick Hall
22 Jan 2017


Machine Learning Overview

What is machine learning?

A field of study that gives computers the ability to learn without being explicitly programmed.

-- Arthur Samuel, 1959





Why now?

- Data, computers, and algorithms are commodities
- Unstructured data
- Increasing competition in business

Estimating a model for inference

What happened? Why?

Assumptions, parsimony,
interpretation

Linear models, statistics

Models tend to be static

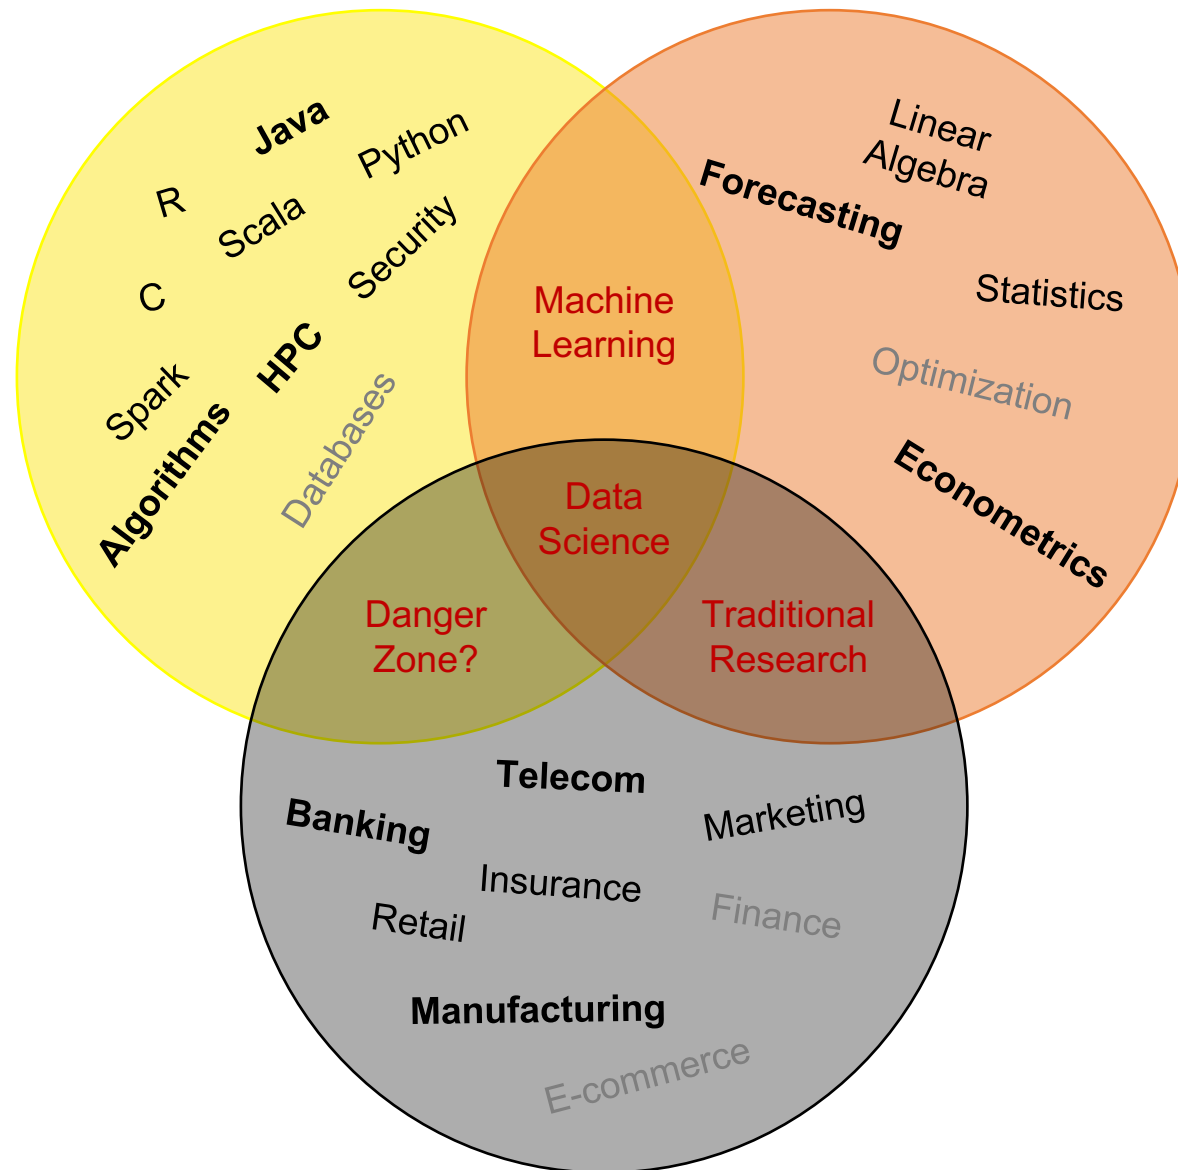
Training a model for prediction

What will happen?

Predictive accuracy,
production deployment

Machine learning

Many models evolve
elegantly



1. There is no perfect language.



If someone claims to have the perfect programming language, he is either a fool or a salesman or both.

-- Bjarne Stroustrup

2. There is no perfect algorithm.



Algorithms that search for an extremum of a cost function perform exactly the same when averaged over all possible cost functions.

-- D.H. Wolpert

3. Doing things right is always hard.



Developing and deploying ML systems is relatively fast and cheap, but maintaining them over time is difficult and expensive.

-- Google,
*Hidden Technical Debt in
Machine Learning Systems*

H₂O.ai Overview

Company Overview


Founded	2011 Venture-backed, debuted in 2012
Products	<ul style="list-style-type: none">• H2O: In-Memory AI Prediction Engine• Sparkling Water: Spark Integration• Steam: Deployment engine• Deep Water: Deep Learning
Mission	Operationalize Data Science, and provide a platform for users to build beautiful data products
Team	70 employees <ul style="list-style-type: none">• Distributed Systems Engineers doing Machine Learning• World-class visualization designers
Headquarters	Mountain View, CA



H₂O.ai

H2O Community Growth

Tremendous Momentum Globally



* DATA FROM GOOGLE ANALYTICS EMBEDDED IN THE END USER PRODUCT

Large User Circle

- 75,000+ users from 8,500+ companies in 142 countries. Top 5 from:

-  [United States](#)
-  [India](#)
-  [Japan](#)
-  [Germany](#)
-  [United Kingdom](#)

H2O.ai Offers AI Open Source Platform

Product Suite to Operationalize Data Science

100% Open Source



In-Memory, Distributed
Machine Learning
Algorithms with Speed
and Accuracy

The logo for Deep Water, featuring the text "Deep Water" in black on a light yellow square background.

State-of-the-art
Deep Learning on GPUs
with TensorFlow, MXNet
or Caffe with the ease of
use of H2O



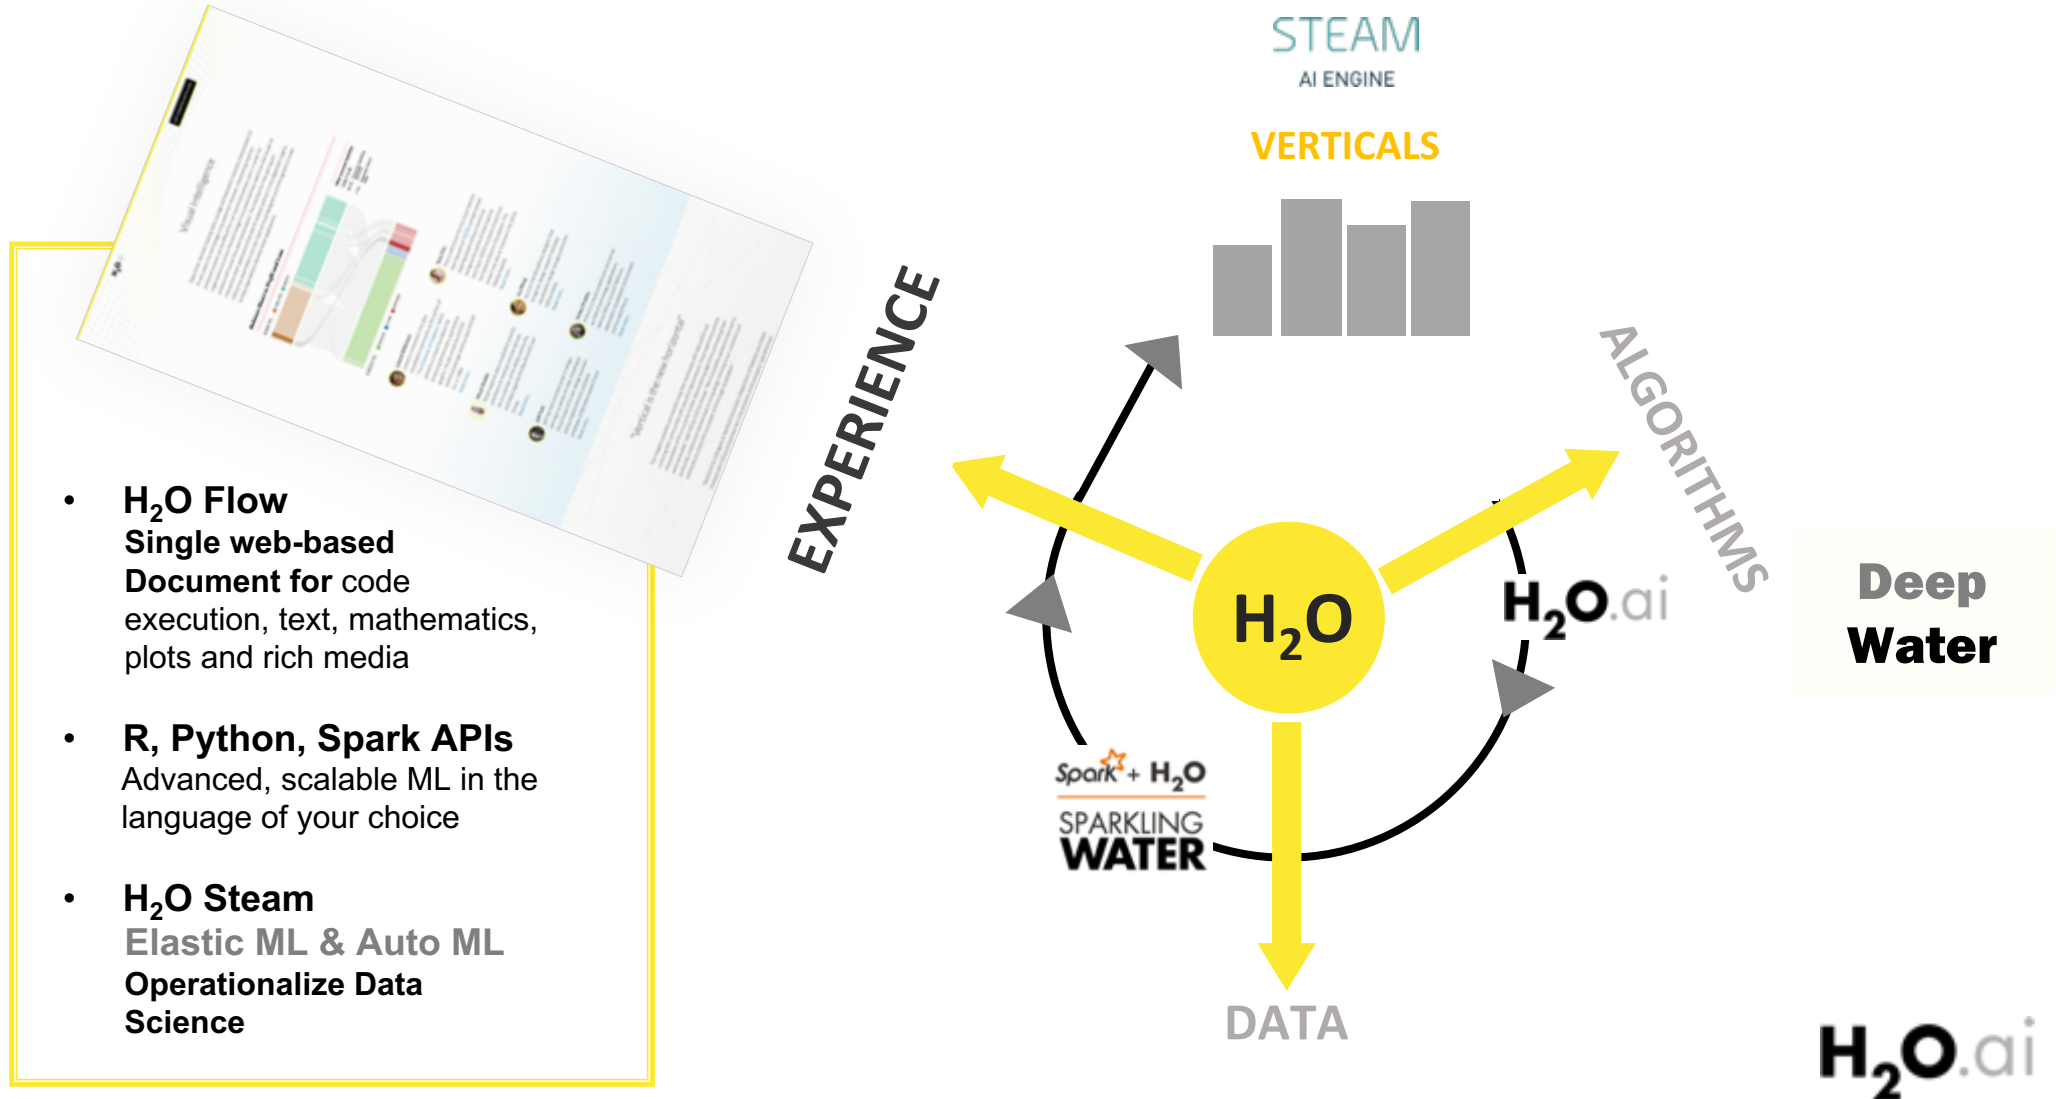
H2O Integration with
Spark. Best Machine
Learning on Spark.

The logo for Steam, featuring the word "Steam" in orange.

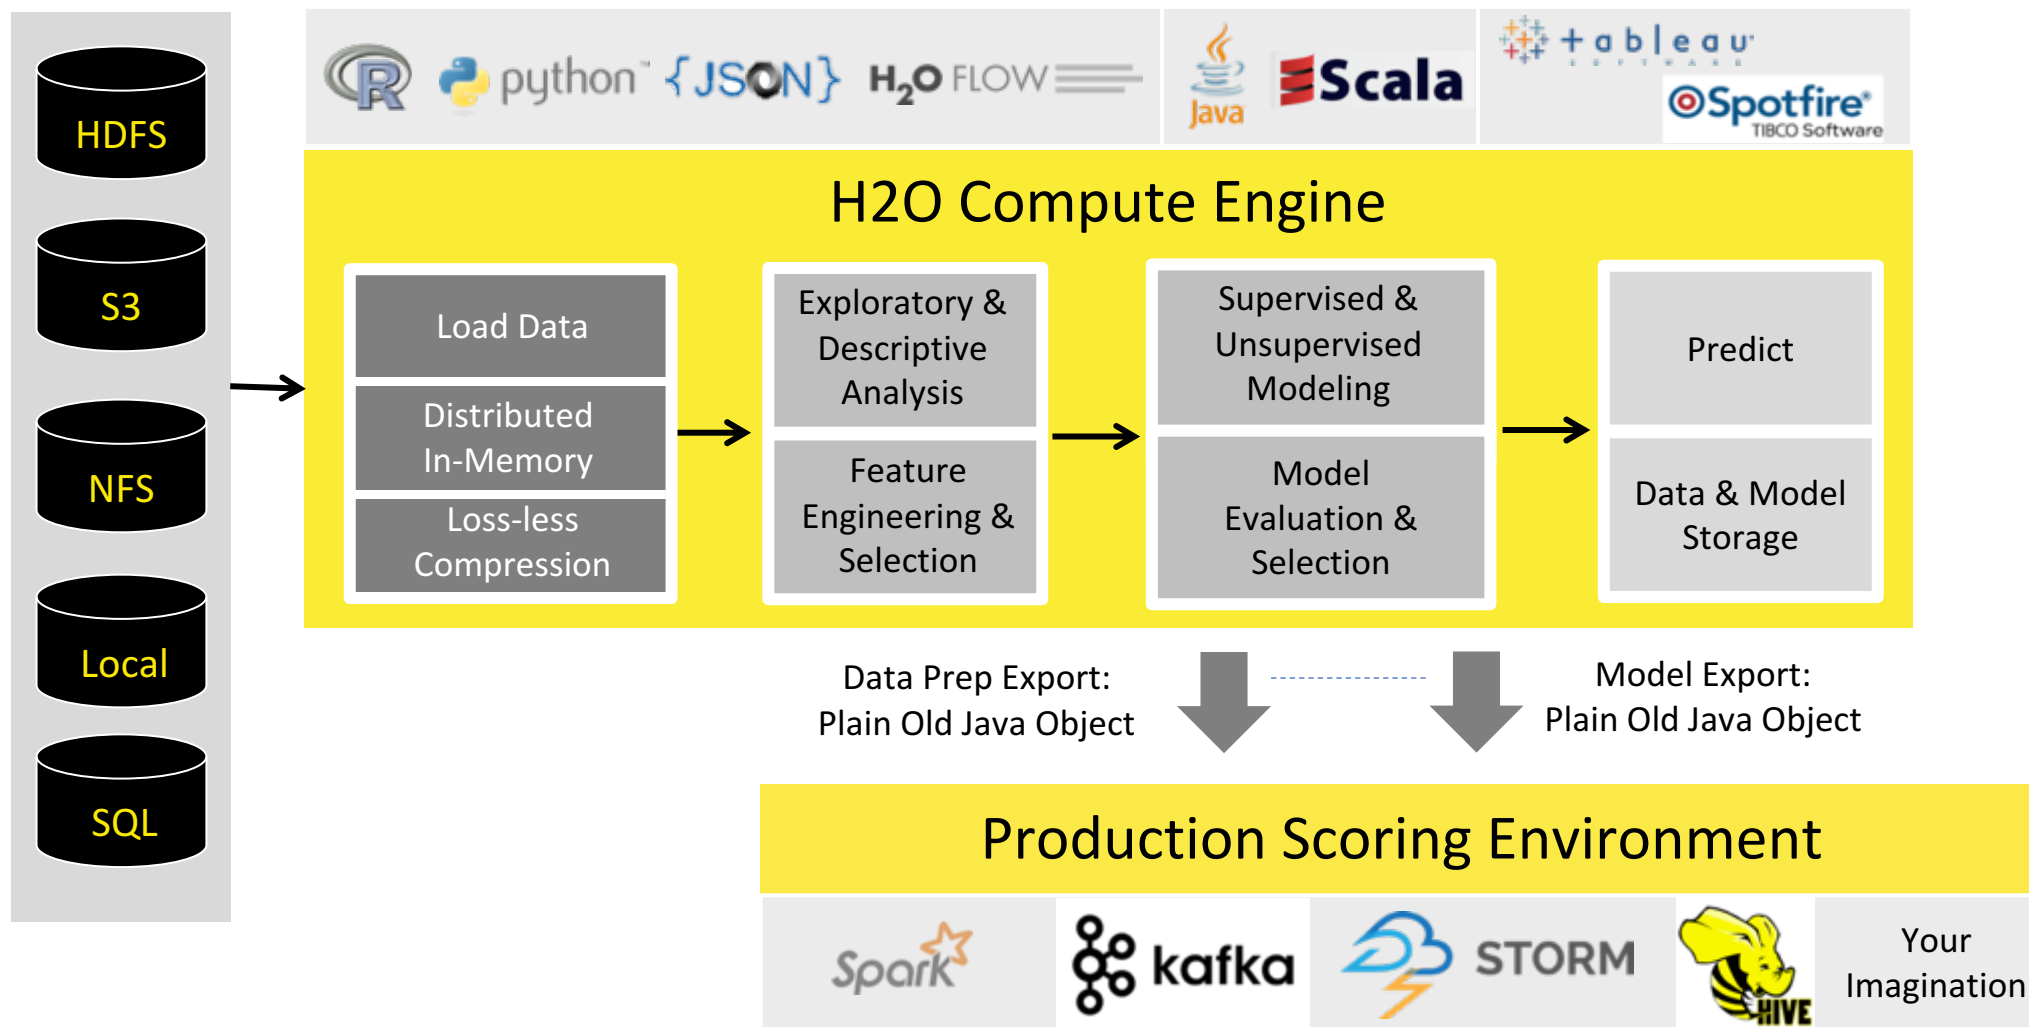
Operationalize and
Streamline Model
Building, Training and
Deployment Automatically
and Elastically

H₂O.ai Now Focused On Experience

Beyond Algorithms and Data



High Level Architecture




Deep Water

Next-Gen Distributed Deep Learning with H₂O

One Interface - GPU Enabled - Significant Performance Gains

Inherits All H₂O Properties in Scalability, Ease of Use and Deployment



H₂O integrates with existing **GPU** backends for **significant performance gains**



Convolutional Neural Networks enabling **Image, video, speech recognition**



Hybrid Neural Network Architectures enabling **speech to text translation, image captioning, scene parsing** and more



Recurrent Neural Networks enabling **natural language processing, sequences, time series**, and more

Bring the Greatest Possible Machine Learning Functionality and Choice to Spark



Spark MLlib

- Elegant APIs
- RDD's
- Multi-tenant Context

H₂O.ai

- Speed and scalability
- Distributed
- Columnar compression
- Fully-featured ML algorithms
- Interactive GUI – Flow
- R, Python, Java, Scala APIs

- Spark 2.0 API compatibility
- Use H2O algorithms in conjunction with, *or instead of*, MLlib algorithms on Spark
- Build Ensembles using H2O and MLlib Algorithms
- Run Spark, MLlib, Scala in Flow
- Export Spark models as POJOs
- Toolchain for ML pipelines and debugging support

Steam

Automated Platform to Build and Scale Smart Data Products

Smart
Data
Product

Smart
Data
Product

Smart
Data
Product

Smart
Data
Product



AI – Machine Learning



Automation



Scalability

H₂O.ai Makes a Difference as an AI Platform

Open Source



- 100% open source

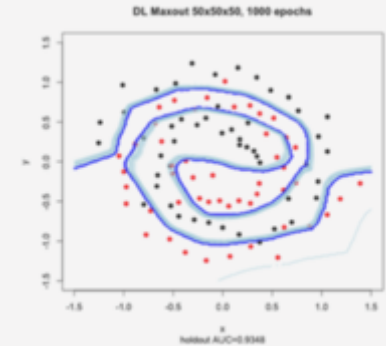
Big Data Ecosystem




Flexible Interface



Smart and Fast Algorithms



Scalability and Performance



- Distributed In-Memory Computing Platform
- Distributed Algorithms
- Fine-Grain MapReduce

Rapid Model Deployment

- Highly portable models deployed in Java (POJO) and Model Object Optimized (MOJO)
- Automated and streamlined scoring service deployment with Rest API



GPU Enablement



Cloud Integration



Questions?